

Problem Set 3

Codename: Jambul

Late Submission: No

1 Problem 1

Algorithm 1. (Chirp Z-transform.) We will use the Fast Fourier Transform for the convolution of two polynomials to solve the problem.

1. Create two empty vectors a, b with $|a| = n$, $|b| = 2n - 1$, and set $a_i = x_i z^{i^2/2}$, $b_i = b_{2n-i-1} = z^{-i^2/2}$.
2. Set $a_i = 0$ for $n < i \leq 2n - 1$, so $|a| = 2n - 1$.
3. Compute the FFT of a and b and multiply the resultant sequences pointwise. Denote the output of these operations by c .
4. Compute the inverse FFT of c , which we denote by $(a * b)$.
5. Return $z^{j^2/2} \cdot (a * b)$.

Proof of correctness. As per the hint, we have that

$$ij - \frac{j^2}{2} = \frac{i^2}{2} - \frac{(j-i)^2}{2} \implies y_j = \sum_{i=0}^{n-1} x_i z^{ij} = z^{j^2/2} \sum_{i=0}^{n-1} x_i z^{i^2/2 - (j-i)^2/2}.$$

Now define $a = (a_i)_i$, $b = (b_i)_i$ to be two new sequences defined by

$$a_i = x_i z^{i^2/2}, b_i = z^{-i^2/2}.$$

Then we can write

$$y_j = z^{j^2/2} \sum_{i=0}^{n-1} a_i b_{j-i}.$$

We now define two new vectors a^*, b^* , with $a_i^* = a_i$ for $0 \leq i \leq n - 1$, and $a_i^* = 0$ for $n - 1 < i \leq 2n - 2$, and with $b_i^* = b_{i+n-1}$. Thus we simply append zeroes to a until its length is $2n - 1$ to get a^* , and we shift the index on b by $n - 1$ to get b^* .¹ Then $|a| = |b| = 2n - 1$, and we have that

$$y_j = z^{j^2/2} \cdot (a * b)_j.$$

We know from lecture that the convolution $(a * b)$ can be obtained by computing the FFT of a and b , multiplying the transforms pointwise, and taking the inverse FFT of the resultant sequence. Thus the algorithm indeed returns the desired chirp z-transform (y_0, \dots, y_{n-1}) of (x_0, \dots, x_{n-1}) .

Complexity. We will now confirm that the overall runtime of this algorithm is $\mathcal{O}(n \lg(n))$. We must first compute the sequences a, b . Computing $z^{i^2/2}$ will take at most $\mathcal{O}(\lg(n))$ time (and multiplication of x_i by the resultant complex number is assumed to take $\mathcal{O}(1)$ time), so computing both a and b takes $\mathcal{O}(n \lg(n))$. Padding a then takes $\mathcal{O}(n)$ time. Finally, computing the FFT of both a and b takes $\mathcal{O}(n \lg(n))$ time; multiplying the resultant transforms pointwise takes $\mathcal{O}(n)$ time; and finding the inverse FFT of the multiplied sequence takes $\mathcal{O}(n \lg(n))$.

Thus the algorithm indeed has $\mathcal{O}(n \lg(n))$ time complexity.

¹Note that in Algorithm 1, the vectors we define in step 1 and 2 are simply a^* and b^* .

2 Problem 2

We will use a combination of a divide-and-conquer approach alongside the FFT for polynomial multiplication to solve the problem.

Algorithm 2. (Polynomial Multiplication.) Let our input array of constants be $c = (c_1, \dots, c_n)$. Then we proceed as follows.

1. If $|c| = 1$, return $(c_1, 1)$.
2. Set $m = \lfloor n/2 \rfloor$.
3. Set l to be the list of coefficients returned from calling “Polynomial Multiplication” on (c_1, \dots, c_m) .
4. Set r to be the list of coefficients returned from calling “Polynomial Multiplication” on (c_{m+1}, \dots, c_n) .
5. Compute the FFTs of l and r and multiply the resultant sequences pointwise.
6. Compute the inverse FFT of the sequence of pointwise products and denote the output o .
7. Return the coefficients on the terms of o in order of decreasing degree.

Proof of correctness. Denote the list of input constants by c , so in any call to Algorithm 2, we are given some partition $c^* = (c_i, \dots, c_j)$, $1 \leq i \leq n$, $1 \leq j \leq n$, of c . We will prove the correctness of the algorithm by induction on the length $|c^*|$ of c^* .

For our base case, we consider $|c^*| = 1$, that is, the event in which our input polynomial is of the form $(x - c_i)$ for some real constant c_i . Trivially, then, the coefficients of this polynomial are given by $(c_i, 1)$. Thus the claim holds for the base case.

Now assume the claim holds for some $1 < |c^*| = k < n$. Then assume we are given some list of constants of length $k + 1$. We will notate this new list again by c^* .² By assumption, calling the algorithm on the left half $(c_1^*, \dots, c_{\lfloor (k+1)/2 \rfloor}^*)$ and the right half $(c_{\lfloor (k+1)/2 \rfloor + 1}^*, \dots, c_{k+1}^*)$ of c^* will yield the coefficients representing the polynomials formed by $(x - c_1^*) \dots (x - c_{\lfloor (k+1)/2 \rfloor}^*)$ and $(x - c_{\lfloor (k+1)/2 \rfloor + 1}^*) \dots (x - c_{k+1}^*)$, respectively.

Then, we know from lecture that the product of two polynomials l and r is given by

$$o(x) = \{l * r\}(x) = \mathcal{F}^{-1}\{\mathcal{F}(l) \cdot \mathcal{F}(r)\}, \quad (1)$$

where $\mathcal{F}\{\cdot\}$ is the Fourier transform operator. It follows that computing and returning the coefficients of (1) will yield the list of coefficients of the polynomial

$$(x - c_1^*) \dots (x - c_{\lfloor n/2 \rfloor}^*) \cdot (x - c_{\lfloor n/2 \rfloor + 1}^*) \dots (x - c_n^*),$$

which is of course c^* . Thus by induction on $|c^*|$, the claim holds for all $1 \leq |c^*| \leq n$, and in particular, it holds for $c^* = c$.

Complexity. As mentioned, the algorithm combines a divide-and-conquer approach with the FFT algorithm for polynomial multiplication to solve the problem. Splitting the polynomial takes $\mathcal{O}(1)$ time. We know we can leverage the FFT to compute the coefficients of the product of two polynomials in $\mathcal{O}(n \log(n))$ time. We then must do this $\mathcal{O}(\log(n))$ times, as we will be “merging” (multiplying polynomials) $\mathcal{O}(\log(n))$ times, since we split the array of constants in half at each call. Thus the total running time of the algorithm is indeed $\mathcal{O}(n \log^2(n))$.

²Apologies for the notation abuse.

3 Problem 3

We will use a variation of the mergesort algorithm to solve this problem.

In particular, we note that we seek to populate the output entries $c[i]$ with the number of elements in a that begin on the right of $a[i]$ but ultimately rest on the left side of it in the sorted array. We will first outline a subroutine for the modified mergesort algorithm that will be the driving component of this solution.

Algorithm 3. (Modified Mergesort.) This algorithm requires some input array $m = (m_1, \dots, m_n)$, where each entry m_i is a 2-tuple with the first entry equal to the index i of the entry (i.e., $m_i[0] = i$). We also assume there exists an array p that has been created outside the algorithm but can be modified by the algorithm.

1. If $|m| = n = 1$: return m .
2. Let m_l be the left half (rounded down) of the input array m ; i.e., $m_{(l)} = (m_1, \dots, m_{\lfloor n/2 \rfloor})$. Let $m_{(r)}$ be the remaining right half of m .
3. Recursively call “Modified Mergesort” on $m_{(l)}$; denote the output s_l for the sorted left half of m .
4. Recursively call “Modified Mergesort” on $m_{(r)}$; denote the output s_r for the sorted right half of m .
5. Create an empty, temporary array o . Set $l = r = 0$.
6. While $l < |s_l|$ and $r < |s_r|$:
 - (a) If $s_l[l][1] < s_r[r][1]$:
 - i. Set the entry at the index $s_l[l][0]$ in the output array p to r .
 - ii. Append $s_l[l]$ to the temporary array o .
 - iii. Add 1 to l .
 - (b) Else:
 - i. Append $s_r[r]$ to the temporary array o .
 - ii. Add 1 to r .
7. While $l < |s_l|$:
 - (a) Set the entry at the index $s_l[l][0]$ in the output array p to r .
 - (b) Append $s_l[l]$ to the temporary array o .
 - (c) Add 1 to l .
8. While $r < |s_r|$:
 - (a) Append $s_r[r]$ to the temporary array o .
 - (b) Add 1 to r .
9. Return the temporary array o .

Algorithm 4. (Smaller Elements Count.) We are now ready to outline the full solution to the problem. In line with the problem, we denote the input array a .

1. Create an array t such that $t[i] = (i, a[i])$. Initialize the output array $p = [0, \dots, 0]$ with length $|a| = n$.
2. Call “Modified Mergesort” on t . (We assume in “Modified Mergesort” that the subroutine has permission to modify p .)

3. Return the output array p .

Proof of correctness. Denote the input array $a = (a_1, \dots, a_n)$. We will show that any call to Algorithm 3 with an input $m = (m_1, \dots, m_n)$, $m_i = (i, a_i)$, $a_i \in \mathbb{R}$ will return m sorted by the values stored in the second index of each 2-tuple m_i and will populate an empty array $p = (p_1, \dots, p_n)$ such that $p_i = |\{a_j : a_j < a_i, j > i\}|$. We will prove the correctness of the algorithm by induction on the length n of the input array a .

For our base case, consider input arrays of length $n = 1$. Trivially, we will have that $m = ((0, a_1))$ is in fact sorted by the second entry in each tuple in m , as there is only one such entry in the first place. Then we expect Algorithm 3 to return m and Algorithm 4 to return $p = (0)$, which is indeed the case, since Algorithm 4 initializes $p = (0, \dots, 0)$ with length $|p| = n = 1$.

Now suppose the claim holds for some $n \geq 1$. Then suppose we input a list $m = ((0, a_1), \dots, (0, a_{n+1}))$ to Algorithm 3 through Algorithm 4. Then we know that when we call Algorithm 3 recursively on l the left half of m and r the right half of m , we will be able to retrieve l and r sorted by their entries in the second slot of each 2-tuple in the arrays. Furthermore, also following from the induction hypothesis, we know $p = (p_1, \dots, p_n)$ will be populated such that each entry p_i on the left half of p will hold the number of elements a_j in the left half of a that are of index greater than a_i and are also less than a_i . The same will hold for the right half of p .

Then to combine and sort the arrays l and r , we can simply use the mergesort algorithm described in lecture. To populate p , however, we note we must keep track of the number n_r of elements that have been added to the auxiliary output array. Each time we append an element l_i from l to the auxiliary array, we know that l_i must be greater than every element that is now to the left of it. However, we cannot simply add the number of these elements to p_i : We must include only those elements that are smaller *and* from r . Thus we add n_r as defined previously.

By assumption, this will yield p such that $p_i = |\{a_j : a_j < a_i, j > i\}|$, since each time we add l_i to r_i , we know that p_i has already been updated to include the number of elements in l that were originally to the right of l_i and are smaller than l_i . Thus adding that same count against the new set of numbers in r completes the algorithm, thus confirming its correctness.

Complexity. Similar to the classic mergesort algorithm, this solution has two parts: Splitting and merging. Splitting is done in constant time. Merging, however, takes $\mathcal{O}(n)$ time, as at most we will encounter n comparisons in each merge. Furthermore, at each step, writing to the output array p in the "Modified Mergesort" subroutine takes constant time. Since the array is split into two at each step, there will be $\mathcal{O}(\lg(n))$ levels in the mergesort tree, resulting in the usual $\mathcal{O}(n \lg(n))$ runtime found in mergesort.